

# Assignment 3

## Part A: Lazy Memory Allocation

- When a process needs extra memory, it indicates this requirement using the **sbrk system call** in the xv6 OS.
- **sbrk** uses **growproc** (defined in `proc.c`) to allocate extra memory by mapping virtual addresses to physical addresses in the page table.
- The assignment's goal is to implement **Lazy Memory Allocation**, where memory is allocated when accessed, not when requested.
- To achieve this, the **growproc** call inside **sbrk** is commented out, giving the process the illusion of allocated memory.
- Upon accessing the "allocated" memory, a **page fault** occurs, generating a **T\_PGFLT trap** to the kernel.
- The trap is handled in **trap.c** by calling **handlePageFault()**, which uses **rcr2()** to get the virtual address of the fault.
- The virtual address is rounded to the nearest page boundary and then a free physical page is allocated using **kalloc** from a list of free pages.
- The physical page is mapped to the virtual address using **mappages**, which is declared by removing its static keyword in `trap.c`.
- **mappages** takes the page table, virtual address, physical memory address (converted using **V2P**), and permission bits.

- **walkpgdir()** is used to find the page table entry for a given virtual address, operating with a two-level page table structure.
- **mappages** checks if the page table entry is already mapped (by checking the PRESENT bit) and maps the physical page if not already mapped.
- If an error occurs (such as a remap), an error is raised by **mappages**.

```

95     case T_PGFLT:
96         if(handlePageFault()<0){
97             cprintf("Could not allocate page. Sorry.\n");
98             panic("trap");
99         }
100     break;

```

(trap.c)

```

19 int handlePageFault(){
20     int addr=rcr2();
21     int rounded_addr = PGROUNDDOWN(addr);
22     char *mem=kalloc();
23     if(mem!=0){
24         memset(mem, 0, PGSIZE);
25         if(mappages(myproc()->pgdir, (char*)rounded_addr, PGSIZE, V2P(mem),
26                     PTE_W|PTE_U)<0)
27             return -1;
28         return 0;
29     } else
30         return -1;
31 }

```

(trap.c)

```

60  int
61  mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
62  {
63      char *a, *last;
64      pte_t *pte;
65
66      a = (char*)PGROUNDDOWN((uint)va);
67      last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
68      for(;;){
69          if((pte = walkpgdir(pgdir, a, 1)) == 0)
70              return -1;
71          if(*pte & PTE_P)
72              panic("remap");
73          *pte = pa | perm | PTE_P;
74          if(a == last)
75              break;
76          a += PGSIZE;
77          pa += PGSIZE;
78      }
79      return 0;
80  }
81

```

(vm.c)

```

34 // create any required page table pages.
35 static pte_t *
36 walkpgdir(pde_t *pgdir, const void *va, int alloc)
37 {
38     pde_t *pde;
39     pte_t *pgtab;
40
41     pde = &pgdir[PDX(va)];
42     if(*pde & PTE_P){
43         pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
44     } else {
45         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
46             return 0;
47         // Make sure all those PTE_P bits are zero.
48         memset(pgtab, 0, PGSIZE);
49         // The permissions here are overly generous, but they can
50         // be further restricted by the permissions in the page table
51         // entries, if necessary.
52         *pde = V2P(pgtab) | PTE_P | PTE_W | PTE_U;
53     }
54     return &pgtab[PTX(va)];
55 }

```

(vm.c)

## PART B : Question and answer

### Q1: How does the kernel know which physical pages are used and unused?

xv6 maintains a linked list of free pages in `kalloc.c` called `kmem`. Initially, the list is empty, so xv6 calls `kinit1` through `main()` which adds 4MB of free pages to the list.

### Q2: What data structures are used to answer this question?

A linked list named `freelist` as shown in the above image. Every node of the linked list is a structure defined in `kalloc.c` namely `struct run` (pages are typecast to `(struct run *)` when inserting into `freelist` in `kfree(char *v)`).

### Q3: Where do these reside?

This linked list is declared inside `kalloc.c` inside a structure `kmem`. Every node is of the type `struct run` which is also defined inside `kalloc.c`.

### Q4: Does xv6 memory mechanism limit the number of user processes?

Due to a limit on the size of `ptable` (a max. of `NPROC` elements which is set to 64 by default), the number of user processes are limited in xv6. `NPROC` is defined in `param.h`.

### Q5: If so, what is the lowest number of processes xv6 can 'have' at the same time (assuming the kernel requires no memory whatsoever)?

The lowest number of processes xv6 can "have" at the same time, assuming the kernel requires no memory at all (a theoretical assumption), would still be **2 processes**. Here's why:

#### 1. Process 0 (the "scheduler" process):

- xv6 always starts with an initial process called **process 0** or the **scheduler process**. This process is responsible for managing CPU scheduling and process switching. It is created during the system initialization in `main()`.
- Even though the scheduler process doesn't run user code, it is necessary for xv6 to function because it coordinates all other processes.

#### 2. Process 1 (the "init" process):

- After initializing the kernel and creating the scheduler process, xv6 creates a special process called **init** (process 1). The **init process** is the first user-level process in xv6 and is responsible for starting the shell and other user-level programs.
- This process also serves as a parent for any orphaned processes (i.e., processes whose parent has exited).



## Part B Assignment.

### Task 1:

- **The `create_kernel_process()` function:** This function is where the kernel process is created.
- **Kernel mode:** The kernel process will always remain in kernel mode, meaning it has direct access to the system's hardware and resources.
- **Trapframe:** The kernel process doesn't need to initialize its trapframe because it won't be switching to user mode. Trapframes store user-space register values, which are only relevant for user processes.
- **User space and page table:** Similarly, the kernel process doesn't need a user space or a user section of its page table since it operates exclusively in kernel mode.
- **EIP register:** The EIP (instruction pointer) register is set to the `entry_point` address. This is the starting address of the code that the kernel process will execute.
- **Allocproc and setupkvm:** These functions are used to allocate a process slot in the process table (ptable) and set up the kernel part of the process's page table, respectively. The kernel page table maps virtual addresses above `KERNBASE` (the base address of the kernel) to physical addresses between 0 and `PHYSTOP` (the end of physical memory).

This code creates the kernel process, sets up its initial context, and configures its memory management. The kernel process will then start executing the code at the specified

entry\_point.

```
482 void create_kernel_process(const char *name, void (*entrypoint)()){
483
484     struct proc *p = allocproc();
485
486     if(p == 0)
487         panic("create_kernel_process failed");
488
489     //Setting up kernel page table using setupkvm
490     if((p->pgdir = setupkvm()) == 0)
491         panic("setupkvm failed");
492
493     //This is a kernel process. Trap frame stores user space registers. We don't need to
494     //initialise tf.
495     //Also, since this doesn't need to have a userspace, we don't need to assign a size to
496     //this process.
497
498     //eip stores address of next instruction to be executed
499     p->context->eip = (uint)entrypoint;
500
501     safestrcpy(p->name, name, sizeof(p->name));
502
503     acquire(&ptable.lock);
504     p->state = RUNNABLE;
505     release(&ptable.lock);
506 }
```

(proc.c)



## Task 2:

This task has various parts. First, we need a process queue that keeps track of the processes that were refused additional memory since there were no free pages available. We created a circular queue struct called `rq`. And the specific queue that holds processes with swap out requests is `rqueue`. We have also created the functions corresponding to `rq`, namely `rpush()` and `rpop()`. The queue needs to be accessed with a lock that we have initialized in `pinit`. We have also initialized the initial values of `s` and `e` to zero in `userinit`. Since the queue and the functions relating to it are needed in other files too, we added prototypes in `defs.h` too.

```
172 struct rq{
173     struct spinlock lock;
174     struct proc* queue[NPROC];
175     int s;
176     int e;
177 };
178
179 //circular request queue for swapping out requests.
180 struct rq rqueue;
181
182 struct proc* rpop(){
183
184     acquire(&rqueue.lock);
185     if(rqueue.s==rqueue.e){
186         release(&rqueue.lock);
187         return 0;
188     }
189     struct proc *p=rqueue.queue[rqueue.s];
190     (rqueue.s)++;
191     (rqueue.s)%=NPROC;
192     release(&rqueue.lock);
193
194     return p;
195 }
```

(proc.c)

```

384 void
385 pinit(void)
386 {
387     initlock(&ptable.lock, "ptable");
388     initlock(&rqueue.lock, "rqueue");
389     initlock(&sleeping_channel_lock, "sleeping_channel");
390     initlock(&rqueue2.lock, "rqueue2");
391 }

```

(proc.c)

```

509 void
510 userinit(void)
511 {
512     acquire(&rqueue.lock);
513     rqueue.s=0;
514     rqueue.e=0;
515     release(&rqueue.lock);
516
517     acquire(&rqueue2.lock);
518     rqueue2.s=0;
519     rqueue2.e=0;
520     release(&rqueue2.lock);
521
522     struct proc *p;
523     extern char _binary_initcode_start[], _binary_initcode_size[];
524
525     p = allocproc();
526
527     initproc = p;
528     if((p->pgdir = setupkvm()) == 0)
529         panic("userinit: out of memory?");
530     inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);
531     p->sz = PGSIZE;
532     memset(p->tf, 0, sizeof(*p->tf));
533     p->tf->cs = (SEG_UCODE << 3) | DPL_USER;
534     p->tf->ds = (SEG_UDATA << 3) | DPL_USER;
535     p->tf->es = p->tf->ds;
536     p->tf->ss = p->tf->ds;
537     p->tf->eflags = FL_IF;
538     p->tf->esp = PGSIZE;
539     p->tf->eip = 0; // beginning of initcode.S
540
541     safestrcpy(p->name, "initcode", sizeof(p->name));
542     p->cwd = namei("/");
543
544     // this assignment to p->state lets other cores
545     // run this process. the acquire forces the above
546     // writes to be visible, and the lock is also needed
547     // because the assignment might not be atomic.
548     acquire(&ptable.lock);
549
550     p->state = RUNNABLE;
551
552     release(&ptable.lock);
553
554 }

```

(proc.c)

```

197 int rpush(struct proc *p){
198
199     acquire(&rqueue.lock);
200     if((rqueue.e+1)%NPROC==rqueue.s){
201         release(&rqueue.lock);
202         return 0;
203     }
204     rqueue.queue[rqueue.e]=p;
205     rqueue.e++;
206     (rqueue.e)%=NPROC;
207     release(&rqueue.lock);
208
209     return 1;
210 }

```

(proc.c)

```

129 extern struct rq rqueue;
130 extern struct rq rqueue2;
131 int rpush(struct proc *p);
132 struct proc* rpop();
133 struct proc* rpop2();
134 int rpush2(struct proc* p);
135

```

(defs.h)

Now, whenever kalloc is not able to allocate pages to a process, it returns zero. This notifies allocvm that the requested memory wasn't allocated (mem=0). Here, we first need to change the process state to sleeping.

(\*Note: The process sleeps on a special sleeping channel called sleeping\_channel that is secured by a lock called sleeping\_channel\_lock. sleeping\_channel\_count is used for corner cases when the system boots)

Then, we need to add the current process to the swap out request queue, rqueue:

```

14 struct spinlock sleeping_channel_lock;
15 int sleeping_channel_count=0;
16 char * sleeping_channel;

```

(vm.c)

```

int
allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
{
    char *mem;
    uint a;

    if(newsz >= KERNBASE)
        return 0;
    if(newsz < oldsz)
        return oldsz;

    a = PGROUNDUP(oldsz);
    for(; a < newsz; a += PGSIZE){
        mem = kalloc();
        if(mem == 0){
            // cprintf("allocuvm out of memory\n");
            deallocuvm(pgdir, newsz, oldsz);

            //SLEEP
            myproc()->state=SLEEPING;
            acquire(&sleeping_channel_lock);
            myproc()->chan=sleeping_channel;
            sleeping_channel_count++;
            release(&sleeping_channel_lock);

            rpush(myproc());
            if(!swap_out_process_exists){
                swap_out_process_exists=1;
                create_kernel_process("swap_out_process", &swap_out_process_function);
            }

            return 0;
        }
        memset(mem, 0, PGSIZE);
        if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
            cprintf("allocuvm out of memory (2)\n");
            deallocuvm(pgdir, newsz, oldsz);
            kfree(mem);
            return 0;
        }
    }
    return newsz;
}

```

(vm.c)

\*Note: create\_kernel\_process here creates a swapping out kernel process to allocate a page for this process if it doesn't already exist. When the swap out process ends, the swap\_out\_process\_exists (declared as extern in defs.h and initialized in proc.c to 0) variable is set to 0. When it is created, it is set to 1 (as seen above). This is done so multiple swap out processes are not created. swap\_out\_process is explained later.

Next, we create a mechanism by which whenever free pages are available, all the processes sleeping on sleeping\_channel are woken up. We edit kfree in kalloc.c in the following way:

Basically, all processes that were preempted due to lack of availability of pages were sent sleeping on the sleeping channel. We wake all processes currently sleeping on sleeping\_channel by calling the wakeup() system call.

```
60 void
61 kfree(char *v)
62 {
63
64     struct run *r;
65     // struct proc *p=myproc();
66
67     if((uint)v % PGSIZE || v < end || V2P(v) >= PHYSTOP){
68         panic("kfree");
69     }
70
71     // Fill with junk to catch dangling refs.
72     // memset(v, 1, PGSIZE);
73     for(int i=0;i<PGSIZE;i++){
74         v[i]=1;
75     }
76
77     if(kmem.use_lock)
78         acquire(&kmem.lock);
79     r = (struct run*)v;
80     r->next = kmem.freelist;
81     kmem.freelist = r;
82     if(kmem.use_lock)
83         release(&kmem.lock);
84
85     //Wake up processes sleeping on sleeping channel.
86     if(kmem.use_lock)
87         acquire(&sleeping_channel_lock);
88     if(sleeping_channel_count){
89         wakeup(sleeping_channel);
90         sleeping_channel_count=0;
91     }
92     if(kmem.use_lock)
93         release(&sleeping_channel_lock);
94
95 }
```

(kalloc.c)

Now, I will explain the swapping out process. The entry point for the swapping out process is swap\_out\_process\_function. Since the function is very long, I have attached two screenshots:

```

244 void swap_out_process_function()
245 {
246     acquire(&rqueue.lock);
247     while(rqueue.s!=rqueue.e){
248         struct proc *p=rpop();
249
250         pde_t* pd = p->pgdir;
251         for(int i=0;i<NPDETRIES;i++){
252
253             //skip page table if accessed. chances are high, not every page table was accessed.
254             if(pd[i]&PTE_A)
255                 continue;
256             //else
257             pte_t *pgtab = (pte_t*)P2V(PTE_ADDR(pd[i]));
258             for(int j=0;j<NPTETRIES;j++){
259
260                 //Skip if found
261                 if((pgtab[j]&PTE_A) || !(pgtab[j]&PTE_P))
262                     continue;
263                 pte_t *pte=(pte_t*)P2V(PTE_ADDR(pgtab[j]));
264
265                 //for file name
266                 int pid=p->pid;
267                 int virt = ((1<<22)*i)+((1<<12)*j);
268
269                 //file name
270                 char c[50];
271                 int_to_string(pid,c);
272                 int x=strlen(c);
273                 c[x]='_';
274                 int_to_string(virt,c+x+1);
275                 safestrcpy(c+strlen(c),".swp",5);
276
277                 // file management
278                 int fd=proc_open(c, O_CREATE | O_RDWR);
279                 if(fd<0){
280                     cprintf("error creating or opening file: %s\n", c);
281                     panic("swap_out_process");
282                 }
283
284                 if(proc_write(fd,(char *)pte, PGSIZE) != PGSIZE){
285                     cprintf("error writing to file: %s\n", c);
286                     panic("swap_out_process");
287                 }
288                 proc_close(fd);
289
290                 kfree((char*)pte);
291                 memset(&pgtab[j],0,sizeof(pgtab[j]));
292
293                 //mark this page as being swapped out.
294                 pgtab[j]=((pgtab[j])^(0x080));
295
296                 break;
297             }
298         }
299     }
300 }
301
302 release(&rqueue.lock);
303
304 struct proc *p;
305 if((p=myproc())==0)
306     panic("swap out process");
307
308 swap_out_process_exists=0;
309 p->parent = 0;
310 p->name[0] = '*';
311 p->killed = 0;
312 p->state = UNUSED;
313 sched();
314 }

```

(proc.c)

The process runs a loop until the swap out requests queue (rqueue1) is non-empty. When the queue is empty, a set of instructions are executed for the termination of swap\_out\_process. The loop starts by popping the first process from rqueue and uses the LRU policy to determine a victim page in its page table. We iterate through each entry in the process' page table (pgdir) and extract the physical address for each secondary page table. For each secondary page table, we iterate through the page table and look at the accessed bit (A) on each of the entries (The accessed bit is the sixth bit from the right. We check if it is set by checking the bitwise & of the entry and PTE\_A (which we defined as 32 in mmu.c)).

Important note regarding the Accessed flag: Whenever the process is being context switched into by the scheduler, all accessed bits are unset. Since we are doing this, the accessed bit seen by swap\_out\_process function will indicate whether the entry was accessed in the last iteration of the process:

```
752     for(int i=0;i<NPENTRIES;i++){
753         //If PDE was accessed
754
755         if(((p->pgdir)[i]&PTE_P && ((p->pgdir)[i]&PTE_A){
756
757             pte_t* pgtab = (pte_t*)P2V(PTE_ADDR((p->pgdir)[i]));
758
759             for(int j=0;j<NPENTRIES;j++){
760                 if(pgtab[j]&PTE_A){
761                     pgtab[j]^=PTE_A;
762                 }
763             }
764
765             ((p->pgdir)[i])^=PTE_A;
766         }
767     }
768
769     // Switch to chosen process. It is the process's job
770     // to release ptable.lock and then reacquire it
771     // before jumping back to us.
772     c->proc = p;
773     switchvm(p);
774     p->state = RUNNING;
775
776     switch(&(c->scheduler), p->context);
777     switchvm();
```

(proc.c)

This code resides in the scheduler and it basically unsets every accessed bit in the process' page table and its secondary page tables.

Now, back to swap\_out\_process function. As soon as the function finds a secondary page table entry with the accessed bit unset, it chooses this entry's physical page number (using macros mentioned in part A report) as the victim page. This page is then swapped out and stored to drive.

We use the process' pid (pid, line 267 in image) and virtual address of the page to be eliminated (virt, line 268 in image) to name the file that stores this page. We have

created a new function called int to\_string that copies an integer into a given string. We use this function to make the filename using integers pid and virt. Here is that function declared in proc.c:

```

149 void int_to_string(int x, char *c){
150     if(x==0)
151     {
152         c[0]='0';
153         c[1]='\0';
154         return;
155     }
156     int i=0;
157     while(x>0){
158         c[i]=x%10+'0';
159         i++;
160         x/=10;
161     }
162     c[i]='\0';
163
164     for(int j=0;j<i/2;j++){
165         char a=c[j];
166         c[j]=c[i-j-1];
167         c[i-j-1]=a;
168     }
169
170 }

```

(proc.c)

We need to write the contents of the victim page to the file with the name <pid>\_<virt>.swp. But we encounter a problem here. We store the filename in a string called c. File system calls cannot be called from proc.c. The solution was that we copied the open, write, read, close etc. functions from sysfile.c to proc.c, modified them since the sysfile.c functions used a different way to take arguments and then renamed them to proc\_open, proc\_read, proc\_write, proc\_close etc. so we can use them in proc.c. Some examples:

```

33 int
34 proc_write(int fd, char *p, int n)
35 {
36     struct file *f;
37     if(fd < 0 || fd >= NOFILE || (f=myproc()->ofile[fd]) == 0)
38         return -1;
39     return filewrite(f, p, n);
40 }
41

```



```

20  int
21  proc_close(int fd)
22  {
23      struct file *f;
24
25      if(fd < 0 || fd >= NOFILE || (f=myproc()->ofile[fd]) == 0)
26          return -1;
27
28      myproc()->ofile[fd] = 0;
29      fileclose(f);
30      return 0;
31  }
32

```

There are many more functions (proc\_open, proc\_fdalloc etc.) and you can check them out in proc.c. I can't paste all of them here.

Now, using these functions, we write back a page to storage. We open a file (using proc\_open) with O\_CREATE and O\_RDWR permissions (we have imported fcntl.h with these macros). O\_CREATE creates this file if it doesn't exist and O\_RDWR refers to read/write. The file descriptor is stored in an integer called fd. Using this file descriptor, we write the page to this file using proc\_write. Then, this page is added to the free page queue using kfree so it is available for use (remember we also wake up all processes sleeping on sleeping\_channel when kfree adds a page to the free queue). We then clear the page table entry too using memset.

After this, we do something important: for Task 3, we need to know if the page that caused a fault was swapped out or not. In order to mark this page as swapped out, we set the 8th bit from the right ( $2^7$ ) in the secondary page table entry. We use xor to accomplish this task (LINE 295 in image).

Suspending kernel process when no requests are left:

When the queue is empty, the loop breaks and suspension of the process is initiated. While exiting the kernel processes that are running, we can't clear their kstack from within the process because after this, they will not know which process to execute next. We need to clear their kstack from outside the process. For this, we first preempt the process and wait for the scheduler to find this process. When the scheduler finds a kernel process in the UNUSED state, it clears this process' kstack and name. The scheduler identifies the kernel process in unused state by checking its name in which the first character was changed to '\*' when the process ended.

The ending of kernel processes has two parts:

1. From within process:

```
303
304     struct proc *p;
305     if((p=myproc())==0)
306         panic("swap out process");
307
308     swap_out_process_exists=0;
309     p->parent = 0;
310     p->name[0] = '*';
311     p->killed = 0;
312     p->state = UNUSED;
313     sched();
314 }
```

(proc.c)

2. From Scheduler

```
name.
741     if(p->state==UNUSED && p->name[0]=='*'){
742
743         kfree(p->kstack);
744         p->kstack=0;
745         p->name[0]=0;
746         p->pid=0;
747     }
```

(proc.c)

## TASK 3

We first needed to create a swap in request queue. We used the same struct (rg) as in Task 2 to create a swap in request queue called rqueue2 in proc.c. We also declare an extern prototype for rqueue2 in defs.h. Along with declaring the queue, we also created the corresponding functions for rqueue2 (rpop2() and rpush2()) in proc.c and declared their prototype in defs.h. We also initialized its lock in pinit. We also initialized its s and e variables in userinit.

Next, we add an additional entry to the **struct proc** in **proc.h** called **addr (int)**. This entry will tell the swapping in function at which virtual address the page fault occurred:

Proc.h(in struct proc):

```
51 char name[16];           // Process name (debugging)
52 int addr;                // ADDED: virtual address of pagefault
```

Next, we need to handle page fault (**T\_PGFLT**) traps raised in trap.c. We do it in a function called handlePageFault():

Trap.c:

```
104 case T_PGFLT:
105     | handlePageFault();
106     break;
107 //PAGEBREAK: 13
```

```

19 void handlePageFault(){
20     int addr=rcr2();
21     struct proc *p=myproc();
22     acquire(&swap_in_lock);
23     sleep(p,&swap_in_lock);
24     pde_t *pde = &(p->pgdir)[PDX(addr)];
25     pte_t *pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
26
27     if((pgtab[PTX(addr)])&0x080){
28         //This means that the page was swapped out.
29         //virtual address for page
30         p->addr = addr;
31         rpush2(p);
32         if(!swap_in_process_exists){
33             swap_in_process_exists=1;
34             create_kernel_process("swap_in_process", &
35                                   swap_in_process_function);
36         }
37     } else {
38         exit();
39     }
}

```

(trap.c)

In handlePageFault, just like Part A, we find the virtual address at which the page fault occurred by using rcr2(). We then put the current process to sleep with a new lock called swap\_in\_lock (initialized in trap.c and with extern in defs.h). We then obtain the page table entry corresponding to this address (the logic is identical to walkpgdir). Now, we need to check whether this page was swapped out. In Task 2, whenever we swapped out a page, we set its page table entry's bit of 7th order ( $2^7$ ). This is mentioned at the beginning of the 5th page of this report. Thus, in order to check whether the page was swapped out or not, we check its 7th order bit using bitwise & with 0x080. If it is set, we initiate swap\_in\_process (if it doesn't already exist - check using swap\_in\_process\_exists). Otherwise, we safely suspend the process using exit() as the assignment asked us to do.

Now, we go through the swapping in process. The entry point for the swapping out process is swap\_in\_process\_function (declared in proc.c) as you can see in handlePageFault.

Note: swap\_in\_process\_function is shown on the next page since it is long. Refer to the next page for the actual function.

I have already mentioned how we have implemented file management functions in `proc.c` in the Task 2 part of the report. I will just mention which functions I used and how I used them here. The function runs a loop until `rqueue2` is not empty. In the loop, it pops a process from the queue and extracts its `pid` and `addr` value to get the file name. Then, it creates the filename in a string called `"c"` using `int to string`. Then, it used `proc_open` to open this file in read only mode (`O_RDONLY`) with file descriptor `fd`. We then allocate a free frame (`mem`) to this process using `kalloc`. We read from the file with the `fd` file descriptor into this free frame using `proc_read`. We then make `mappages` available to `proc.c` by removing the static keyword from it in `vm.c` and then declaring a prototype in `proc.c`. We then use `mappages` to map the page corresponding to `addr` with the physical page that got using `kalloc`; and read into (`mem`). Then we wake up, the process for which we allocated a new page to fix the page fault using `wakeup`. Once the loop is completed, we run the kernel process termination instructions.

```
18 int mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm);  
(proc.c)
```

```

325 void swap_in_process_function(){
326
327     acquire(&rqueue2.lock);
328     while(rqueue2.s!=rqueue2.e){
329         struct proc *p=rpop2();
330
331         int pid=p->pid;
332         int virt=PTE_ADDR(p->addr);
333
334         char c[50];
335         int_to_string(pid,c);
336         int x=strlen(c);
337         c[x]='_';
338         int_to_string(virt,c+x+1);
339         safestrcpy(c+strlen(c),".swp",5);
340
341         int fd=proc_open(c,O_RDONLY);
342         if(fd<0){
343             release(&rqueue2.lock);
344             cprintf("could not find page file in memory: %s\n", c);
345             panic("swap_in_process");
346         }
347         char *mem=kalloc();
348         proc_read(fd,PGSIZE,mem);
349
350         if(mappages(p->pgdir, (void *)virt, PGSIZE, V2P(mem), PTE_W|PTE_U)<0){
351             release(&rqueue2.lock);
352             panic("mappages");
353         }
354         wakeup(p);
355     }
356
357     release(&rqueue2.lock);
358     struct proc *p;
359     if((p=myproc())==0)
360         panic("swap_in_process");
361
362     swap_in_process_exists=0;
363     p->parent = 0;
364     p->name[0] = '*';
365     p->killed = 0;
366     p->state = UNUSED;
367     sched();
368 }
369

```

## TASK 4: Sanity Test

In this part, our aim is to create a testing mechanism in order to test the functionalities created by us in the previous parts. We will implement a user-space program named memtest that will do this job for us. The implementation of memtest is given below.

```
C memtest.c > ...
1  #include "types.h"
2  #include "stat.h"
3  #include "user.h"
4
5  int compute(int val){
6  return val*val - 4*val + 1;
7  }
8
9  int
10 main(int argc, char* argv[]){
11
12 for(int idx=0; idx<20; idx++){
13 if(!fork()){
14 printf(1, "child number: %d\n", idx+1);
15 printf(1, "Iter | Matched Bytes | Different Bytes\n");
16 printf(1, "-----\n\n");
17
18 for(int iteration=0; iteration<10; iteration++){
19 int *buffer = malloc(4096);
20 for(int index=0; index<1024; index++){
21 buffer[index] = compute(index);
22 }
23
24 int countMatch = 0;
25 for(int index=0; index<1024; index++){
26 if(buffer[index] == compute(index))
27 countMatch += 4;
28 }
29
30 if(iteration < 9)
31 printf(1, " %d | %dB | %dB\n", iteration+1, countMatch,
32 4096-countMatch);
33 else
34 printf(1, " %d | %dB | %dB\n", iteration+1, countMatch,
35 4096-countMatch);
36 }
37 printf(1, "\n");
38 exit();
39 }
40 while(wait() != -1);
41 exit();
42 }
43
```

We can make the following observations by looking at the implementation:

- The main process creates 20 child processes using fork() system call.
- Each child process executes a loop with 10 iterations
- At each iteration, 4096B (4KB) of memory is being allocated using malloc()
- The value stored at index  $i$  of the array is given by the mathematical expression  $i^2 - 4i + 1$  which is computed using compute().

- A counter named countMatch is maintained which stores the number of bytes that contain the right values. This is done by checking the value stored at every index with the value returned by the function for that index.

In order to run memtest, we need to include it in the Makefile under UPROGS and EXTRA to make it accessible to the xv6 user.

On running memtest, we obtain the following output.

```

$ memtest
child number: 1
Iter | Matched Bytes | Different Bytes
-----
1 | 4096B | 0B
2 | 4096B | 0B
3 | 4096B | 0B
4 | 4096B | 0B
5 | 4096B | 0B
6 | 4096B | 0B
7 | 4096B | 0B
8 | 4096B | 0B
9 | 4096B | 0B
10 | 4096B | 0B

child number: 2
Iter | Matched Bytes | Different Bytes
-----
1 | 4096B | 0B
2 | 4096B | 0B
3 | 4096B | 0B
4 | 4096B | 0B
5 | 4096B | 0B
6 | 4096B | 0B
7 | 4096B | 0B
8 | 4096B | 0B
9 | 4096B | 0B
10 | 4096B | 0B

child number: 3
Iter | Matched Bytes | Different Bytes
-----
1 | 4096B | 0B
2 | 4096B | 0B
3 | 4096B | 0B
4 | 4096B | 0B
5 | 4096B | 0B
6 | 4096B | 0B
7 | 4096B | 0B
8 | 4096B | 0B
9 | 4096B | 0B
10 | 4096B | 0B

child number: 4
Iter | Matched Bytes | Different Bytes
-----
1 | 4096B | 0B
2 | 4096B | 0B
3 | 4096B | 0B
4 | 4096B | 0B
5 | 4096B | 0B
6 | 4096B | 0B
7 | 4096B | 0B
8 | 4096B | 0B
9 | 4096B | 0B
10 | 4096B | 0B

child number: 5
Iter | Matched Bytes | Different Bytes
-----
1 | 4096B | 0B
2 | 4096B | 0B
3 | 4096B | 0B
4 | 4096B | 0B
5 | 4096B | 0B
6 | 4096B | 0B
7 | 4096B | 0B
8 | 4096B | 0B
9 | 4096B | 0B
10 | 4096B | 0B

```

As can be seen in the output, our implementation passes the sanity test as all the indices store the correct value.