




L27- FRAME ALLOCATION


Overview of Memory Management

- ❖ Demand Paging
- ❖ Copy-on-Write
- ❖ Page Replacement
- ❖ Allocation of Frames
- ❖ Thrashing


Steps in Handling a Page Fault



Page Replacement



Free Frames Allocation



Before allocation

After allocation

Allocation of Frames

- ❖ **Frame-allocation algorithm** determines
 - ❖ How many frames to give each process?
 - ❖ Which frames to replace?
- ❖ Two major allocation schemes
 - ❖ fixed allocation
 - ❖ priority allocation

Fixed Allocation

- ❖ Equal allocation – For example, if there are 100 frames (after allocating frames for the OS) and 5 processes, give each process 20 frames
 - ❖ Keep some as free frame buffer pool

Proportional Allocation

- ❖ Proportional allocation – Allocate according to the size of process

s_i = size of process p_i

$$S = \sum s_i$$

m = total number of frames

$$a_i = \text{allocation for } p_i = \frac{s_i}{S} \times m$$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 4$$

$$a_2 = \frac{127}{137} \times 64 \approx 57$$

Priority Allocation

- ❖ Use a proportional allocation scheme using priorities rather than size
- ❖ If process P_i generates a page fault,
 - ❖ select for replacement one of its frames
 - ❖ select for replacement a frame from a process with lower priority number


Global vs. Local Allocation

- ❖ **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another
 - ❖ But then process execution time can vary greatly
 - ❖ But greater throughput so more common
- ❖ **Local replacement** – each process selects from only its own set of allocated frames
 - ❖ More consistent per-process performance
 - ❖ But possibly underutilized memory

Thrashing

- ❖ If a process does not have enough pages, the page-fault rate is high
 - ❖ Page fault to get page
 - ❖ Replace existing frame
 - ❖ But quickly need replaced frame back
 - ❖ This leads to:
 - ❖ Low CPU utilization
 - ❖ Operating system thinking that it needs to increase the degree of multiprogramming
 - ❖ Another process added to the system
 - ❖ **Thrashing** \equiv a process is busy swapping pages in and out

Thrashing




Demand Paging and Thrashing

- ❖ Why does demand paging work? - **Locality model**
 - ❖ Process migrates from one locality to another
 - ❖ Localities may overlap
- ❖ Why does thrashing occur?
 - ❖ Σ size of locality > total memory size
 - ❖ Limit effects by using local or priority page replacement

Working-Set Model

- ❖ $\Delta \equiv$ working-set window \equiv a fixed number of page references
Example: 10,000 instructions
- ❖ WS_i (working set of Process P_i) = Pages referenced in the most recent Δ
 - ❖ if Δ too small will not encompass entire locality
 - ❖ if Δ too large will encompass several localities
 - ❖ if $\Delta = \infty \Rightarrow$ will encompass entire program



- ❖ $D = \sum WS_i \equiv$ total demand frames (Approximation of locality)
- ❖ if $D > m \Rightarrow$ Thrashing; if $D > m$, then suspend/swap out processes

How to compute Working-Set?


- ❖ Approximate with interval timer + a reference bit
- ❖ Example: $\Delta = 10,000$: Timer interrupts after every 5000 time units
 - ❖ 2 history bits for each page is kept in memory
 - ❖ Whenever a timer interrupts, copy the reference bit to history bit.
 - ❖ Sets the values of all reference bits to 0
 - ❖ During page fault, if one of the history bits = 1 \Rightarrow page in working set

How to compute Working-Set?

- ❖ Why counter, history and reference bits approach not completely accurate?
- ❖ Improvement = 10 bits and interrupt every 1000 time units


Page-Fault Frequency

- ❖ More direct approach than WSS
- ❖ Establish acceptable **page-fault frequency (PFF)** rate and use local replacement policy
 - ❖ If actual rate too low, process loses frame
 - ❖ If actual rate too high, process gains frame



Working Sets and Page Fault Rates

- ❖ Direct relationship between working set of a process and its page-fault rate
- ❖ Working set changes over time
- ❖ Peaks and valleys over time

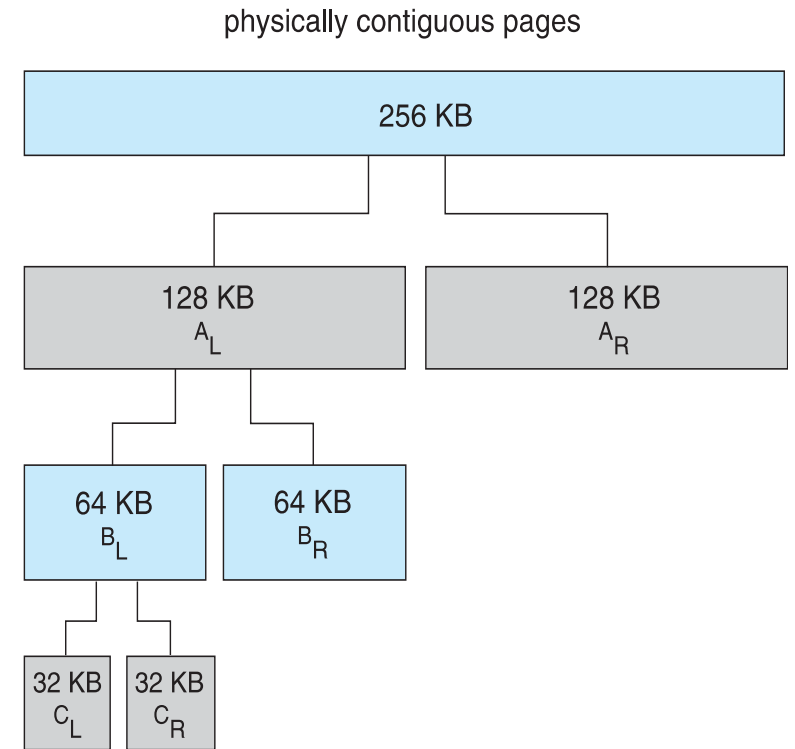


Buddy System

- ❖ Allocates memory from fixed-size segment consisting of physically-contiguous pages
- ❖ Memory allocated using power-of-2 allocator
 - ❖ Satisfies requests in units sized as power of 2
 - ❖ Request rounded up to next highest power of 2
 - ❖ When smaller allocation needed than is available, current chunk split into two buddies of next-lower power of 2
 - ❖ Continue until appropriate sized chunk available
- ❖ Advantage – quickly coalesce unused chunks into larger chunk
- ❖ Disadvantage - fragmentation

Buddy System Allocator

- ❖ Assume 256KB chunk available
- ❖ Kernel requests 21KB
 - ❖ Split into A_L and A_R of 128KB each
 - ❖ One further divided into B_L and B_R of 64KB
 - ❖ One further into C_L and C_R of 32KB each
 - ❖ One used to satisfy request



Prepaging

- ❖ Reduce the large number of page faults that occurs at process startup
- ❖ Prepage all or some of the pages a process will need, before they are referenced
- ❖ But if prepaged pages are unused, I/O and memory was wasted
- ❖ Assume s pages are prepaged and α of the pages is used
- ❖ Cost of $s * \alpha$ saved pages faults vs cost of prepaging $s * (1 - \alpha)$ unnecessary pages
- ❖ α near zero \Rightarrow prepaging loses

Page Size

- ❖ Sometimes OS designers have a choice on custom-built CPU
- ❖ Page size selection criteria:
 - ❖ Fragmentation and Resolution
 - ❖ Page table size
 - ❖ I/O overhead
 - ❖ Number of page faults
 - ❖ Locality
 - ❖ TLB size and effectiveness
- ❖ Always power of 2, usually in the range 2^{12} to 2^{22}

TLB Reach

- ❖ TLB Reach - The amount of memory accessible from the TLB
- ❖ $\text{TLB Reach} = (\text{TLB Size}) \times (\text{Page Size})$
- ❖ Ideally, the working set of each process is stored in the TLB
 - ❖ Otherwise there is more time spend in resolving memory references in page table (delay).
- ❖ Increase the Page Size
 - ❖ This may lead to an increase in fragmentation as not all applications require a large page size
- ❖ Provide Multiple Page Sizes
 - ❖ This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation

Program Structure

- ❖ `int[128,128] data;` Each row is stored in one page
- ❖ A page can store 128 words

Program 1 [128 x 128 = 16,384 page faults]

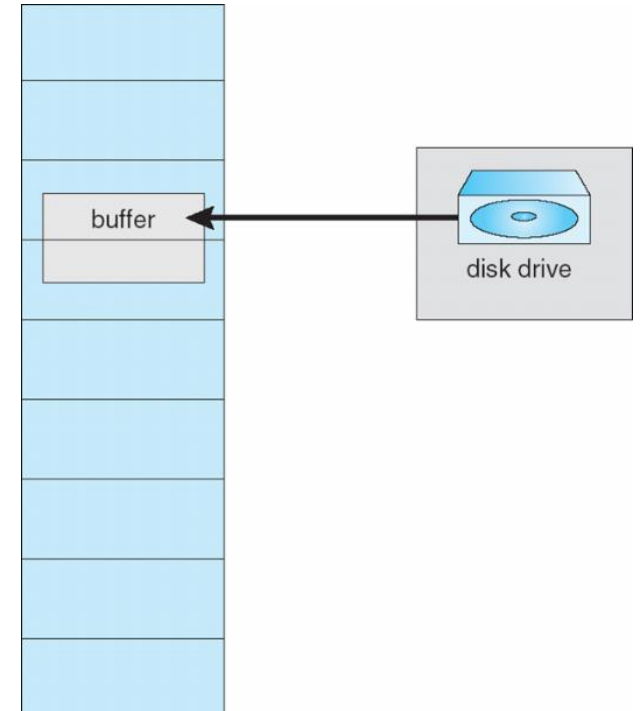
```
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i, j] = 0;
```

Program 2 [128 page faults]

```
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        data[i, j] = 0;
```


I/O interlock

- ❖ **I/O Interlock** – Pages must sometimes be locked into memory
- ❖ Consider I/O - Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm
- ❖ When I/O is complete pages are unlocked





Thank You