



# LECTURE 5: PROCESS API (SYSTEM CALL)

## PROCESS CREATION

Chapters 5 OSTEP:  
[https://pages.cs.wisc.edu/~remzi/OSTEP  
P/cpu-api.pdf](https://pages.cs.wisc.edu/~remzi/OSTEP/cpu-api.pdf)

# AGENDA


- HOW TO CREATE AND CONTROL PROCESSES?
- What interfaces should the OS present for process creation and control? How should these interfaces be designed to enable powerful functionality, ease of use, and high performance

# Operating systems APIs

- This group of APIs defines how applications use the resources and services of operating systems.
- Every OS has its set of APIs, for instance, [Windows API](#) or Linux API (kernel user-space API and [kernel internal API](#)).

# API

- API provided by OS is a set of “system calls”
  - System call is a function call into OS code that runs at a higher privilege level of the CPU
  - Sensitive operations (e.g., access to hardware) are allowed only at a higher privilege level
  - Some “blocking” system calls cause the process to be blocked and descheduled (e.g., read from disk)




# POSIX API

- We want to avoid writing programs for each OS
- POSIX API: a standard set of system calls that an OS must implement
  - Programs written to the POSIX API can run on any POSIX compliant OS
  - Most modern OSes are POSIX compliant
  - Ensures program portability
- Program language libraries hide the details of invoking system calls
  - The printf function in the C library calls the write system call to write to screen
  - User programs usually do not need to worry about invoking system calls

# System calls for process creations and manipulations

- Unix based
- fork() creates a new child process
  - All processes are created by forking from a parent
  - The init process is ancestor of all processes
- exec() makes a process execute a given executable
- exit() terminates a process
- wait() causes a parent to block until child terminates
- Many variants exist of the above system calls with different arguments



# fork()

- System call **fork()** is used to create processes.
- It takes no arguments and returns a process ID.
- The purpose of **fork()** is to create a **new** process, which becomes the *child* process of the caller.
- After a new child process is created, **both** processes will execute the next instruction following the **fork()** system call.
- Therefore, we have to distinguish the parent from the child. This can be done by testing the returned value of **fork()**:
  - If **fork()** returns a negative value, the creation of a child process was unsuccessful.
  - **fork()** returns a zero to the newly created child process.
  - **fork()** returns a positive value, the **process ID** of the child process, to the parent. The returned process ID is of type **pid\_t** defined in **sys/types.h**. Normally, the process ID is an integer. Moreover, a process can use function **getpid()** to retrieve the process ID assigned to this process.
- Therefore, after the system call to **fork()**, a simple test can tell which process is the child. **Please note that Unix will make an exact copy of the parent's address space and give it to the child. Therefore, the parent and child processes have separate address spaces.**

*fork() system call return a process ID.*

*int k = fork();  
k < 0 child creation failed  
k == 0 child created*

*type  
=> pid\_t*

*Parent → child has separate address space.  
Address space*

# Fork Program

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4
5  int main(int argc, char *argv[]) {
6      printf("hello (pid:%d)\n", (int) getpid());
7      int rc = fork();
8      if (rc < 0) {
9          // fork failed
10         fprintf(stderr, "fork failed\n");
11         exit(1);
12     } else if (rc == 0) {
13         // child (new process)
14         printf("child (pid:%d)\n", (int) getpid());
15     } else {
16         // parent goes down this path (main)
17         printf("parent of %d (pid:%d)\n",
18             rc, (int) getpid());
19     }
20     return 0;
21 }
```

```
prompt> ./p1
hello (pid:29146)
parent of 29147 (pid:29146)
child (pid:29147)
prompt>
```



# wait() System Call

- `wait()` System call is used for a parent to wait for a child process to finish what it has been doing.
- Process termination scenarios
  - By calling `exit()` (exit is called automatically when end of main is reached)
  - OS terminates a misbehaving process
- Terminated process exists as a zombie *→ terminated process exists as a ZOMBIE*
- When a parent calls `wait()`, zombie child is cleaned up or “reaped”
- `wait()` blocks in parent until child terminates (non-blocking ways to invoke wait exist)
- What if parent terminates before child?
  - `init` process adopts orphans and reaps them

# fork() and wait()


```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <sys/wait.h>
5
6  int main(int argc, char *argv[]) {
7      printf("hello (pid:%d)\n", (int) getpid());
8      int rc = fork();
9      if (rc < 0) {                // fork failed; exit
10         fprintf(stderr, "fork failed\n");
11         exit(1);
12     } else if (rc == 0) { // child (new process)
13         printf("child (pid:%d)\n", (int) getpid());
14     } else {                    // parent goes down this path
15         int rc_wait = wait(NULL);
16         printf("parent of %d (rc_wait:%d) (pid:%d)\n",
17                rc, rc_wait, (int) getpid());
18     }
19     return 0;
20 }
```

## exec() System Call



- The exec() system call is useful when you want to run a program that is different from the calling program.

- Calling fork() only makes the child executing the same code, not useful.



- A process can run exec() to load another executable to its memory image

- So, a child can run a different program from parent
  - The exec() family of system calls allows a child to break free from its similarity to its parent and execute an entirely new program.

## Variants of exec():

- Functions in the exec() family have different behaviours:
  - $l$  : arguments are passed as a list of strings to the main()
  - $v$  : arguments are passed as an array of strings to the main()
  - $p$  : path/s to search for the new running program
  - $e$  : the environment can be specified by the caller

↪ environment can be specified by the caller.

# Mixing them

- You can mix them, therefore you have:

- `int execl(const char *path, const char *arg, ...);`
- `int execlp(const char *file, const char *arg, ...);`
- `int execlxe(const char *path, const char *arg, ..., char * const envp[]);`
- `int execv(const char *path, char *const argv[]);`
- `int execvp(const char *file, char *const argv[]);`
- `int execvpe(const char *file, char *const argv[], char *const envp[]);`

*It will search automatically.*

- For all of them the initial argument is the name of a file that is to be executed.

# fork(), wait(), and exec()

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <string.h>
5  #include <sys/wait.h>
6
7  int main(int argc, char *argv[]) {
8      printf("hello (pid:%d)\n", (int) getpid());
9      int rc = fork();
10     if (rc < 0) {                // fork failed; exit
11         fprintf(stderr, "fork failed\n");
12         exit(1);
13     } else if (rc == 0) { // child (new process)
14         printf("child (pid:%d)\n", (int) getpid());
15         char *myargs[3];
16         myargs[0] = strdup("wc"); // program: "wc"
17         myargs[1] = strdup("p3.c"); // arg: input file
18         myargs[2] = NULL;          // mark end of array
19         execvp(myargs[0], myargs); // runs word count
20         printf("this shouldn't print out");
21     } else {                      // parent goes down this path
22         int rc_wait = wait(NULL);
23         printf("parent of %d (rc_wait:%d) (pid:%d)\n",
24               rc, rc_wait, (int) getpid());
25     }
26     return 0;
27 }
```

```
prompt> ./p3
hello (pid:29383)
child (pid:29384)
      29      107      1030 p3.c
parent of 29384 (rc_wait:29384) (pid:29383)
prompt>
```

# Case study: How does a shell work?

- In a basic OS, the init process is created after initialization of hardware
- The *init* process spawns a shell like *bash*
- Shell reads user command, forks a child, execs the command executable, waits for it to finish, and reads next command
- Common commands like *ls* are all executables that are simply exec'ed by the shell


prompt>ls

a.txt

b.txt

c.txt


first process is the  
init process  
↓ forks  
Shell → this is  
the shell  
what we  
see, when  
we boot on  
qemu.





# More funky things about the shell

- Shell can manipulate the child in strange ways
- Suppose you want to redirect output from a command to a file
- `prompt>ls > foo.txt`
- Shell spawns a child, rewires its standard output to a file, then calls exec on the child



# All Of The Above With Redirection

```
prompt> ./p4
prompt> cat p4.output
      32      109      846 p4.c
prompt>
```

*executing* →

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <string.h>
5  #include <fcntl.h>
6  #include <sys/wait.h>
7
8  int main(int argc, char *argv[]) {
9      int rc = fork();
10     if (rc < 0) {
11         // fork failed
12         fprintf(stderr, "fork failed\n");
13         exit(1);
14     } else if (rc == 0) {
15         // child: redirect standard output to a file
16         close(STDOUT_FILENO);
17         open("./p4.output", O_CREAT|O_WRONLY|O_TRUNC,
18             S_IRWXU); redirecting
19         // now exec "wc"...
20         char *myargs[3];
21         myargs[0] = strdup("wc"); // program: wc
22         myargs[1] = strdup("p4.c"); // arg: file to count
23         myargs[2] = NULL; // mark end of array
24         execvp(myargs[0], myargs); // runs word count
25     } else {
26         // parent goes down this path (main)
27         int rc_wait = wait(NULL);
28     }
29     return 0;
30 }
```

# Process Control and Users

- Beyond `fork()`, `exec()`, and `wait()`, there are a lot of other interfaces for interacting with processes in UNIX systems
- `kill()` system call : used to send signals to a process
- Dont think that `kill()` is to terminate a process only. It can send all kinds of signals.

*kill() used to Send Signal*



Thank You