



## LECTURE 6: SYSTEM CALL FOR PROCESS MANAGEMENT IN XV6

xv6 book:Pg 7-10  
And refer to xv6 code

# Process System Calls: Shell

- When xv6 boots up, it starts init process (first user process)
- Init forks shell (another user process, which prompts for input)
- Shell executes user commands as follows
  - Shell reads command from terminal
  - Shell forks child (new process created in ptable)
  - When child runs, it calls exec (rewrite code/data with that of command)
  - Shell (parent) waits for child to terminate
  - The whole process repeats again
- Some commands have to be executed by parent process itself, and not by child.
  - For example, “cd” command should change the current directory of parent (shell), not of child
  - Such commands are directly executed by shell itself without forking a child

# Main Function of Shell

```
8700 int
8701 main(void)
8702 {
8703     static char buf[100];
8704     int fd;
8705
8706     // Ensure that three file descriptors are open.
8707     while((fd = open("console", O_RDWR)) >= 0){
8708         if(fd >= 3){
8709             close(fd);
8710             break;
8711         }
8712     }
8713
8714     // Read and run input commands.
8715     while(getcmd(buf, sizeof(buf)) >= 0){
8716         if(buf[0] == 'c' && buf[1] == 'd' && buf[2] == ' '){
8717             // Chdir must be called by the parent, not the child.
8718             buf[strlen(buf)-1] = 0; // chop \n
8719             if(chdir(buf+3) < 0)
8720                 printf(2, "cannot cd %s\n", buf+3);
8721             continue;
8722         }
8723         if(fork1() == 0)
8724             runcmd(parsecmd(buf));
8725         wait();
8726     }
8727     exit();
8728 }
```

→ buf is a string.

→ using "cd" to change directory

changing  
the directory

## What happens on a system call? (1)

- System calls available to user programs are defined in user library header “user.h”
    - Equivalent to C library headers (xv6 doesn't use standard C library)
    - Note that this user code is not available in the PDF source code (which covers only kernel code)
- function which user can use to invoke system call are define in user.h.* →

```
struct stat;
struct rtcdate;

// system calls
int fork(void);
int exit(void) __attribute__((noreturn));
int wait(void);
int pipe(int*);
int write(int, const void*, int);
int read(int, void*, int);
int close(int);
int kill(int);
int exec(char*, char**);
int open(const char*, int);
int mknod(const char*, short, short);
int unlink(const char*);
int fstat(int fd, struct stat*);
int link(const char*, const char*);
int mkdir(const char*);
int chdir(const char*);
int dup(int);
int getpid(void);
char* sbrk(int);
int sleep(int);
int uptime(void);
```



# What happens on a system call? (2)

- System call implementation invokes special “trap” instruction called “int” in x86 (see `usys.S`)  
*↳ “interrupt” in assembly code.*
- The trap (int) instruction causes a jump to kernel code that handles the system call
  - System call number moved into `eax`, to let kernel run the suitable code
  - More on trap instruction later

*int 0x80*

```
#include "syscalls.h"
#include "trap.h"
```

```
#define SYSCALL(name) \
    .globl name; \
    name: \
        movl $SYS_ ## name, %eax; \
        int $T_SYSCALL; \
        ret
```

```
SYSCALL(fork)
SYSCALL(exit)
SYSCALL(wait)
```

*trap instruction (int)*

*causes jump to kernel code.*



*system call's number is stored in "eax" register.*

The term “int” can have different meanings depending on the context in which it is used.

1. “int” as a Trap Instruction:
  - In assembly language or low-level programming, particularly in x86 architecture, “int” stands for “interrupt.” It’s an instruction used to generate a software interrupt. For example, “int 0x80” is commonly used in Linux assembly programs to make system calls. The “int” instruction is used to transfer control to an interrupt handler, which is part of the operating system.
2. “int” as a Data Type (Integer):
  - In high-level programming languages like C, C++, Java, and others, “int” stands for “integer.” It is a data type that represents whole numbers, both positive and negative, including zero. The size of an “int” (in terms of bytes) can vary depending on the programming language and the architecture of the machine.

To sum up, “int” in the context of a trap instruction refers to an “interrupt,” whereas “int” as a data type refers to an “integer” number. The key difference lies in the context—whether you’re dealing with low-level assembly language or high-level programming language data types.

# Fork system call: overview



- Parent allocates new process in ptable, copies parent state to child
- Child process set to runnable, scheduler runs it at a later time

Return value in parent is PID of child, return value in child is set to 0

(Nice)

```

2579 int
2580 fork(void)
2581 {
2582     int i, pid;
2583     struct proc *np;
2584     struct proc *curproc = myproc();
2585
2586     // Allocate process.
2587     if((np = allocproc()) == 0){
2588         return -1;
2589     }
2590
2591     // Copy process state from proc.
2592     if((np->pgdir = copyuvm(curproc->pgdir, curproc->sz)) == 0){
2593         kfree(np->kstack);
2594         np->kstack = 0;
2595         np->state = UNUSED;
2596         return -1;
2597     }
2598     np->sz = curproc->sz;
2599     np->parent = curproc;

```

*new process*

*Page directory (Memory image)*

*Register name can be same.*

*parent has its %eax register.*

*child also has its %eax register.*

```

2600     *np->tf = *curproc->tf;
2601     // Clear %eax so that fork returns 0 in the child.
2602     np->tf->eax = 0;
2603     for(i = 0; i < NOFILE; i++){
2604         if(curproc->ofile[i]){
2605             np->ofile[i] = filedup(curproc->ofile[i]);
2606             np->cwd = idup(curproc->cwd);
2607             safestrcpy(np->name, curproc->name, sizeof(curproc->name));
2608             pid = np->pid;
2609             acquire(&ptable.lock);
2610             np->state = RUNNABLE;
2611             release(&ptable.lock);
2612             return pid;
2613         }
2614     }

```

*child's %eax register.*

*%eax contains process ID of child.*

*if we won't do it "0" then since it is copy of parent & parent's %eax has child's PID.*

*this child's %eax will have its own ID.*

*returning pid to the parent function*

# Exec system call: overview

- Key steps:
  - Copy new executable into memory
  - Create new stack, heap
  - Switch process page table to use new memory image
  - Process begins to run new code after system call ends
- See page 66 of source code PDF for full implementation

# Exit system call: overview

*close  
file* → *child to init* → *mark  
self as zombie* → *Sched()*

```
2626 void
2627 exit(void)
2628 {
2629     struct proc *curproc = myproc();
2630     struct proc *p;
2631     int fd;           ↪ iteration
2632
2633     if(curproc == initproc)
2634         panic("init exiting");
2635
2636     // Close all open files.
2637     for(fd = 0; fd < NOFILE; fd++){
2638         if(curproc->ofile[fd]){
2639             fileclose(curproc->ofile[fd]);
2640             curproc->ofile[fd] = 0;
2641         }
2642     }
2643
2644     begin_op();
2645     iput(curproc->cwd);
2646     end_op();
2647     curproc->cwd = 0;
2648
2649     acquire(&ptable.lock);
```

- Exiting process cleans up state (e.g., close files)
- Pass abandoned children (orphans) to init
- Mark itself as zombie and invoke scheduler

```
2650     // Parent might be sleeping in wait().
2651     wakeup1(curproc->parent);
2652
2653     // Pass abandoned children to init.
2654     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2655         if(p->parent == curproc){
2656             p->parent = initproc;
2657             if(p->state == ZOMBIE)
2658                 wakeup1(initproc);
2659         }
2660     }
2661
2662     // Jump into the scheduler, never to return.
2663     curproc->state = ZOMBIE;
2664     sched();
2665     panic("zombie exit");
2666 }
```



# Wait system call: overview

```
2670 int
2671 wait(void)
2672 {
2673     struct proc *p;
2674     int havekids, pid;
2675     struct proc *curproc = myproc();
2676
2677     acquire(&ptable.lock);
2678     for(;;){
2679         // Scan through table looking for exited children.
2680         havekids = 0;
2681         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2682             if(p->parent != curproc)
2683                 continue;
2684             havekids = 1;
2685             if(p->state == ZOMBIE){
2686                 // Found one.
2687                 pid = p->pid;
2688                 kfree(p->kstack);
2689                 p->kstack = 0;
2690                 freevm(p->pgdir);
2691                 p->pid = 0;
2692                 p->parent = 0;
2693                 p->name[0] = 0;
2694                 p->killed = 0;
2695                 p->state = UNUSED;
2696                 release(&ptable.lock);
2697                 return pid;
2698             }
2699         }
```

```
2700     // No point waiting if we don't have any child
2701     if(!havekids || curproc->killed){
2702         release(&ptable.lock);
2703         return -1;
2704     }
2705
2706     // Wait for children to exit. (See wakeup1 call
2707     sleep(curproc, &ptable.lock);
2708 }
2709 }
```

*if no child  
Sleep.*

*clean up dead  
children found*

Search for dead children in process table

- If dead child found, clean up memory of zombie, return PID of dead child
- If no dead child, sleep until one dies

# Summary of process management system calls in xv6

- Fork – process marks new child's struct proc as RUNNABLE, initializes child memory image and other state that is needed to run when scheduled
- Exec – process reinitializes memory image of user code, data, stack, heap and returns to run new code
- Exit – process marks itself as ZOMBIE, cleans up some of its state, and invokes scheduler
- Wait – parent finds any ZOMBIE child and cleans up all its state. If no dead child yet, it sleeps (marks itself as SLEEPING and invokes scheduler)



Thank You