



LECTURE 13: SCHEDULING ALGO-2

Chapter 3 of Operating System
Concepts- Abraham Silberschatz,
Peter Baer Galvin, Greg Gagne

Chapter 7 of OSTEP

Session Outline

- ❖ **CPU Scheduling**
- ❖ **Limitations of Priority Scheduling**
- ❖ **Priority Inversion**
- ❖ **Multilevel Feedback Queue Scheduling**
- ❖ **Lottery Scheduling**
- ❖ **Summary of Scheduling algorithms**

Scheduling Criteria

- ❖ Different CPU-scheduling algorithms have different properties.
- ❖ Certain characteristics/criteria are used for comparing various CPU scheduling algorithms.
 - ❖ CPU Utilization
 - ❖ Throughput
 - ❖ Turnaround time
 - ❖ Waiting Time
 - ❖ Response Time

CPU Scheduling Algorithms

Batch Systems

- ❖ First-come first-served
- ❖ Shortest job first
- ❖ Shortest remaining Time next

Interactive Systems

- ❖ Round-robin scheduling
- ❖ Priority scheduling
- ❖ Multiple queues
- ❖ Shortest process next
- ❖ Guaranteed scheduling
- ❖ Lottery scheduling
- ❖ Fair-share scheduling

Priority Scheduling

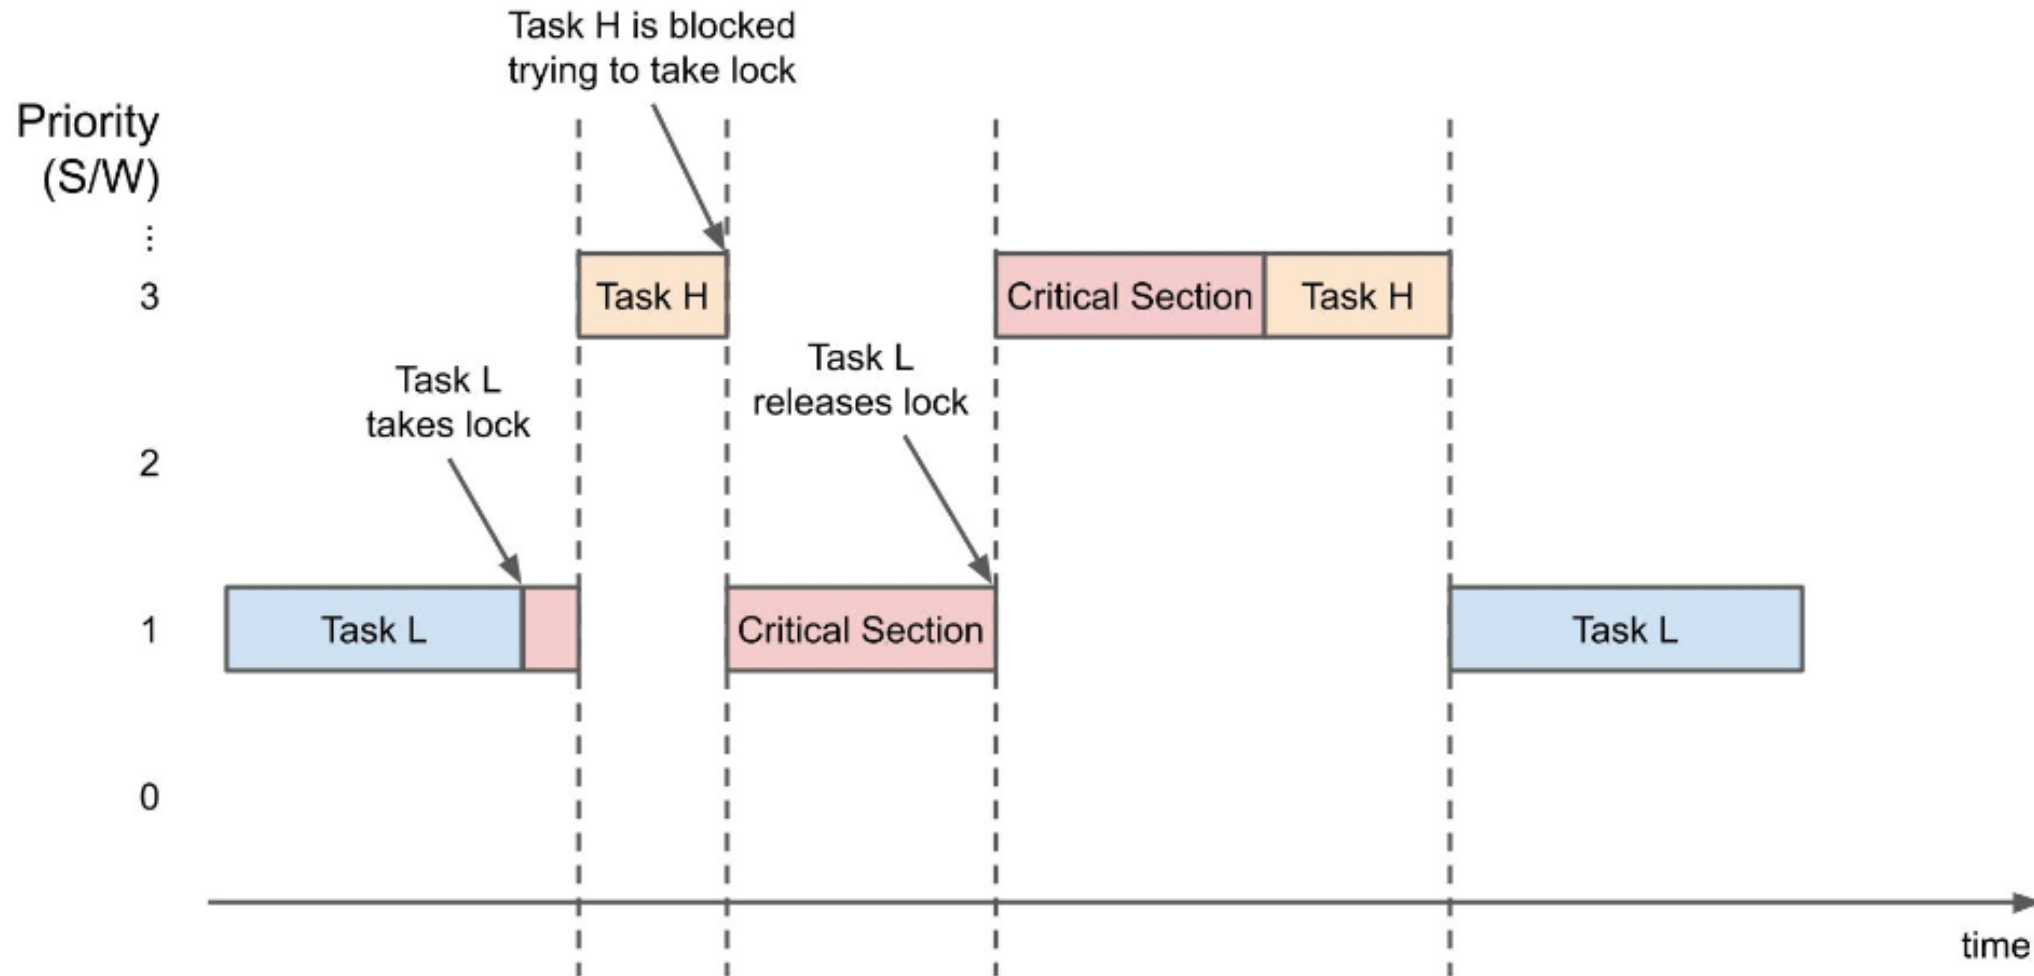
- ❖ Each process has a priority number (integer)
- ❖ Lower the integer, higher the priority
- ❖ Highest priority process is scheduled first; if equal priorities, then FCFS
- ❖ It can have 2 variants; non-preemptive and preemptive
- ❖ Arrival of a new process with a higher priority can preempt the currently running process.

Issues with Priority Scheduling

❖ Priority Inversion

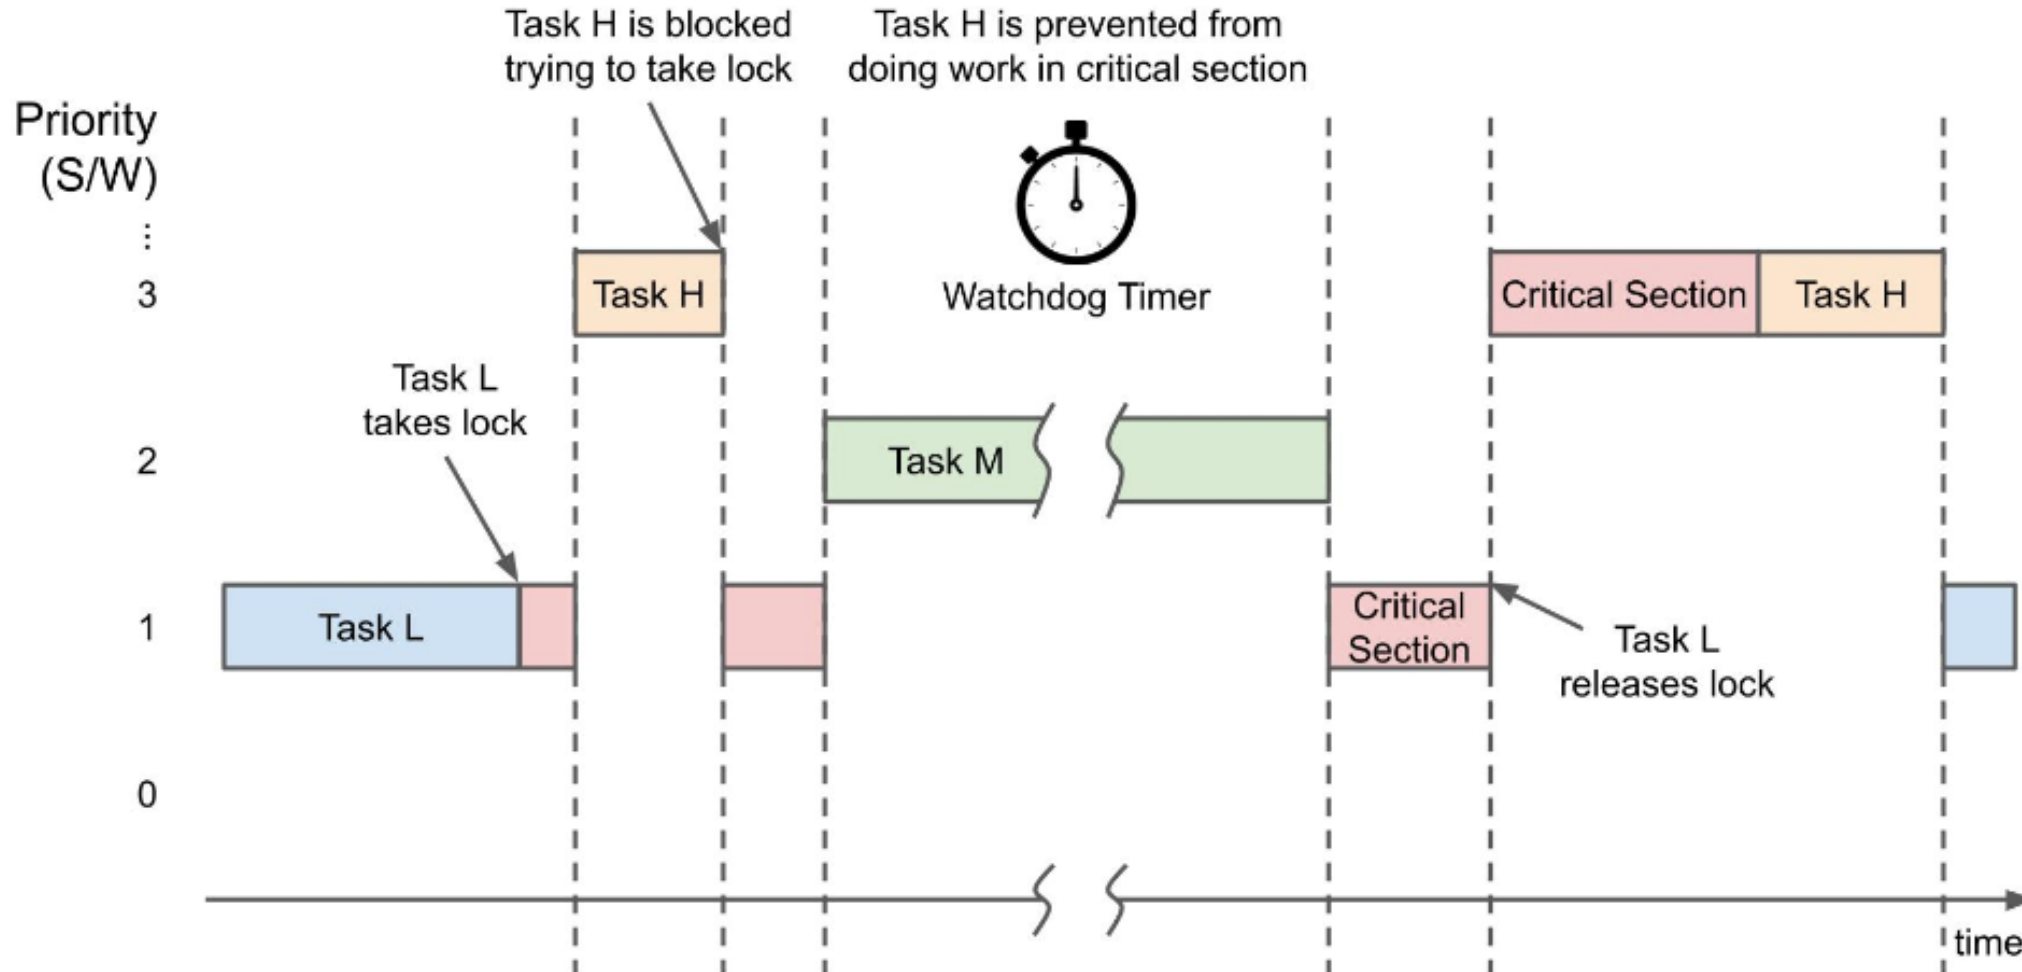
- ❖ Priority inversion is a bug that occurs when a high priority task is indirectly preempted by a low priority task.
- ❖ For example, the low priority task holds a mutex that the high priority task must wait for to continue executing.

Bounded Priority Inversion



As you can see in the diagram above, Task H is blocked so long as Task L holds the lock. The priority of the tasks have been indirectly “inverted” as now Task L is running before Task H.

Unbounded Priority Inversion



Unbounded priority inversion occurs when a medium priority task (Task M) interrupts Task L while it holds the lock. It's called "unbounded" because Task M can now effectively block Task H for any amount of time, as Task M is preempting Task L (which still holds the lock).

Solution To Unbounded Priority Inversion

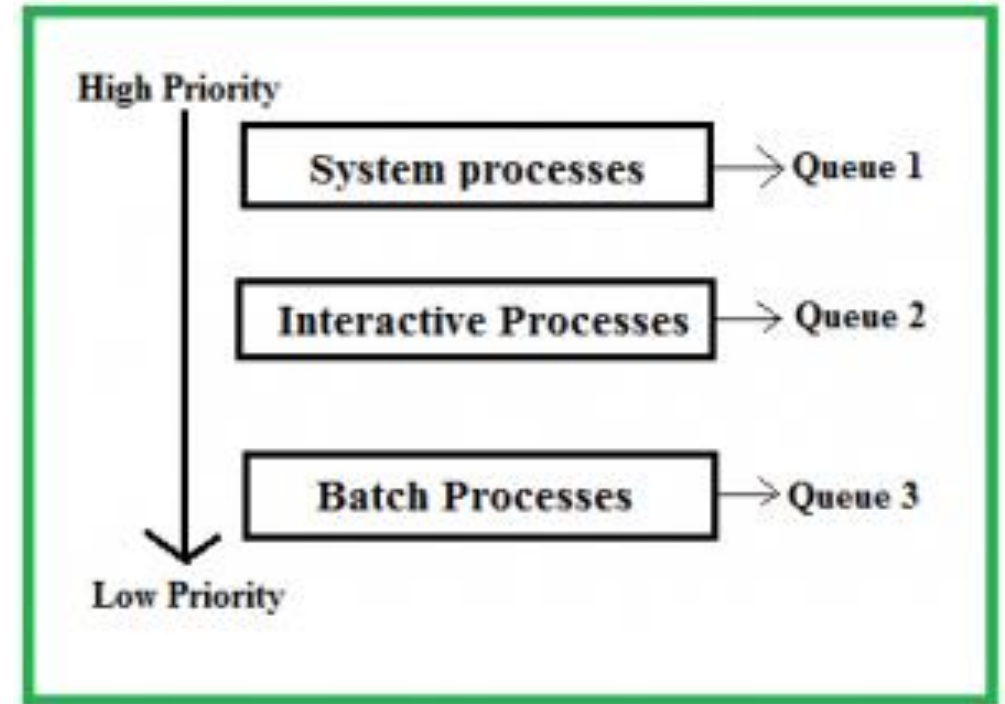
- ❖ Process L should, in fact, be temporarily of higher priority than process M, on behalf of process H.
- ❖ Process H can donate its priority to process L, which, in this case, would make it higher priority than process M.
- ❖ This enables process L to preempt process M and run.
- ❖ When process L is finished, process H becomes unblocked.
- ❖ Process H, now being the highest priority ready process, runs, and process M must wait until it is finished.

Multilevel Queue

- ❖ Ready queue is partitioned into separate queues:
 - ❖ Foreground, interactive process → RR scheduling
 - ❖ Background, batch process → FCFS scheduling
- ❖ A process is permanently assigned to one queue
- ❖ Each queue has its own scheduling algorithm
- ❖ Can be preemptive

Multilevel Feedback Queue Scheduling

- ❖ Scheduling must be done between the queues.
 - ❖ Fixed priority scheduling
 - ❖ Serve all from foreground then from background
 - ❖ Possibility of starvation
- ❖ Time slice
 - ❖ Each queue gets a certain amount of CPU time which it can schedule among its processes
 - ❖ i.e.: 80% Vs 20%




Multilevel Feedback Queue Scheduling

- ❖ **A process can move between the various queues. (Aging)**
- ❖ Multilevel-feedback-queue scheduler defined by the following parameters:
 - ❖ number of queues
 - ❖ scheduling algorithms for each queue
 - ❖ method used to determine when to upgrade a process
 - ❖ method used to determine when to demote a process
 - ❖ method used to determine which queue a process will enter when that process needs service

Example of Multilevel Feedback Queue

- ❖ Three queues:

- ❖ Q0 – time quantum 8 milliseconds, FCFS
- ❖ Q1 – time quantum 16 milliseconds, FCFS
- ❖ Q2 – FCFS



- ❖ A new job enters queue Q0 which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q1.
- ❖ At Q1 job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q2.

Lottery Scheduling

- ❖ Each job some number of lottery tickets are issued
- ❖ On each time slice, randomly pick a winning ticket
- ❖ On average, CPU time is proportional to number of tickets given to each job over time
- ❖ How to assign tickets?
 - ❖ To approximate SRTF, short-running jobs get more, long running jobs get fewer
 - ❖ To avoid starvation, every job gets at least one ticket (everyone makes progress)


Example: Lottery Scheduling

- ❖ Assume short jobs get 10 tickets, long jobs get 1 ticket

# short jobs / # long jobs	% of CPU each short job gets	% of CPU each long job gets
1/1	91%	9%
0/2	N/A	50%
2/0	50%	N/A
10/1	9.9%	0.99%
1/10	50%	5%

Conclusion

- ❖ Scheduling: selecting a waiting process from the ready queue and allocating the CPU to it
- ❖ When do the details of the scheduling policy and fairness really matter?
 - ❖ When there aren't enough resources to go around



Conclusion

- ❖ FCFS scheduling, FIFO Run Until Done:
 - ❖ Simple, but short jobs get stuck behind long ones
- ❖ RR scheduling:
 - ❖ Give each thread a small amount of CPU time when it executes, and cycle between all ready threads
 - ❖ Better for short jobs, but poor when jobs are the same length
- ❖ SJF/SRTF:
 - ❖ Run whatever job has the least amount of computation to do / least amount of remaining computation to do
 - ❖ Optimal (average response time), but unfair; hard to predict the future

Conclusion

- ❖ Multi-Level Feedback Scheduling:
 - ❖ Multiple queues of different priorities
 - ❖ Automatic promotion/demotion of process priority to approximate SJF/SRTF
- ❖ Lottery Scheduling:
 - ❖ Give each thread a number of tickets (short tasks get more)
 - ❖ Every thread gets tickets to ensure forward progress / fairness
- ❖ Priority Scheduling:
 - ❖ Preemptive or Non-preemptive
 - ❖ Priority Inversion



Thank You